

Music Taste Analysis

Ever been asked what sort of music you like and felt unable to describe it convincingly? This notebook represents my effort to once and for all answer the question, because, yes, I regard it to be this complicated.

How to Use

My first pass at this depended upon [Watsonbox's Exportify](#), but I decided I didn't like his version because of bugs and inadequate output detail. So I [went and forked it](#), cleaned up [the code](#), and [hosted it](#) myself.

As such, the code here depends on `.csv` inputs in the format output by [my version](#).

1. To get started, hop on over there, sign in to Spotify to give the app access to your playlists, and export whatever you like.
2. Next, either download this `.ipynb` file and run the notebook yourself or [launch it in Binder](#).
3. Either put the downloaded `.csv` in the same directory as the notebook, or upload it in Binder.
4. Open the `.ipynb` through your browser, update the `filename` variable in the first code cell to point to your playlist instead, and `shift+enter` in each following code cell to generate the corresponding plot. (Or select `Cell -> Run All` from the menu to make all graphs at once.)

Read the Data

For years I've been accumulating my favorite songs in a single master playlist called `music that tickles my fancy`. It's thousands of songs. This is what I'll be analyzing. Let's take a look at the first few rows to get a sense of what we're dealing with.

```
In [1]: filename = 'music_that_tickles_my_fancy.csv'

from matplotlib import pyplot
import seaborn
import pandas
from collections import defaultdict
from scipy.stats import pareto, gamma
```

```
from datetime import date

# read the data
data = pandas.read_csv(filename)
print("total songs:", data.shape[0])
print(data[:3])
```

total songs: 3088

	Spotify ID	Artist IDs \
0	3T9HSgS5jBFdXIBPav51gj	0nJvyjVTb8sAULPYyA1bqU,5yxyJsFanEAuwSM5k0uZKc
1	2bdZDXDoFLzazaomjzoER8	1P6U1dCeHxPui5pIrGmndZ
2	1fE3ddAlmjJ99IIIfLgZjTy	0id62QV2SZZfvBn9xpmuCl

	Track Name \
0	Fanfare for the Common Man
1	Highschool Lover
2	I Need a Dollar

	Album Name \
0	Copland Conducts Copland – Expanded Edition (F...
1	Virgin Suicides
2	I Need A Dollar

	Artist Name(s)	Release Date	Duration (ms) \
0	Aaron Copland,London Symphony Orchestra	1963	196466
1	Air	2000	162093
2	Aloe Blacc	2010-03-16	244373

	Popularity	Added By	Added At	...	Key	Loudness
0	41	spotify:user:pvlkmrv	2014-12-28T00:57:17Z	...	10	-15.727
1	0	spotify:user:pvlkmrv	2014-12-28T00:59:35Z	...	1	-15.025
2	3	spotify:user:pvlkmrv	2014-12-28T01:03:38Z	...	8	-11.825

	Mode	Speechiness	Acousticness	Instrumentalness	Liveness	Valence \
0	1	0.0382	0.986	0.954	0.0575	0.0378
1	0	0.0302	0.952	0.959	0.2520	0.0558
2	0	0.0384	0.178	0.000	0.0863	0.9620

	Tempo	Time Signature
0	104.304	4
1	130.052	4
2	95.516	4

[3 rows x 23 columns]

Artist Bar Chart

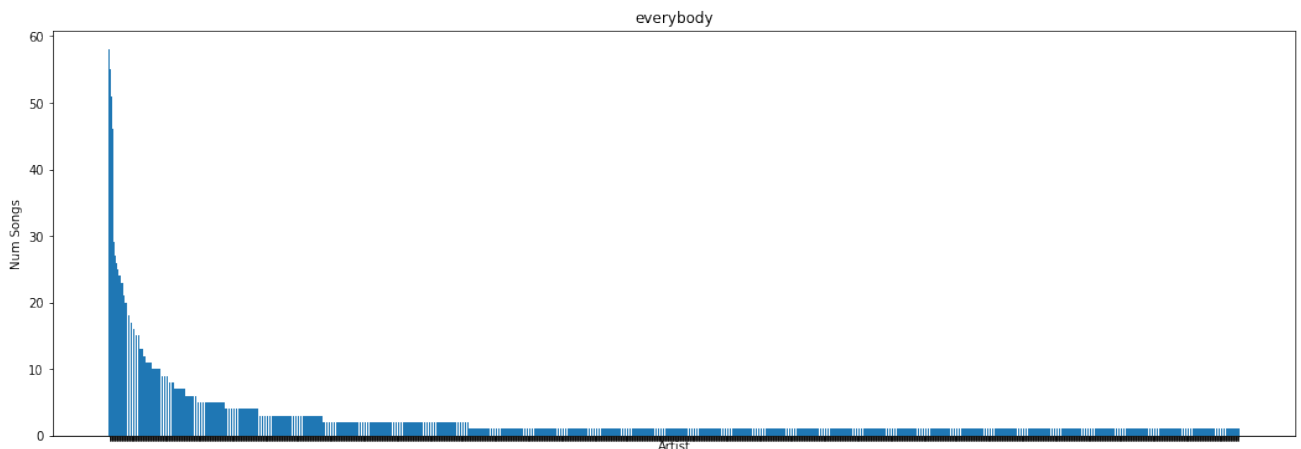
Number of songs binned by artist.

```
In [2]: # count songs per artist
artists = defaultdict(int)
for i,song in data.iterrows():
    for musician in song['Artist Name(s)'].split(','):
        artists[musician] += 1

# sort for chart
artists = pandas.DataFrame(artists.items(), columns=['Artist', 'Num Songs']
                           ).sort_values('Num Songs', ascending=False).reset_
print("number of unique artists:", artists.shape[0])

pyplot.figure(figsize=(18, 6))
pyplot.bar(artists['Artist'], artists['Num Songs'])
pyplot.xticks(visible=False)
pyplot.xlabel(artists.columns[0])
pyplot.ylabel(artists.columns[1])
pyplot.title('everybody')
pyplot.show()
```

number of unique artists: 1387



Note I've attributed songs with multiple artists to multiple bars, so the integral here is the number of unique song-artist pairs, not just the number of songs.

It seems to follow a Pareto distribution. Let's try to fit one.

```
In [3]: # Let's find the best parameters. Need x, y data 'sampled' from the distribu
# parameter fit.
y = []
for i in range(artists.shape[0]):
    for j in range(artists['Num Songs'][i]):
        y.append(i) # just let y have index[artist] repeated for each song

# sanity check. If the dataframe isn't sorted properly, y isn't either.
#pyplot.figure()
#pyplot.hist(y, bins=30)
```

```
# The documentation is pretty bad, but this is okay:
# https://stackoverflow.com/questions/6620471/fitting-empirical-distribution
# ones-with-scipy-python
param = pareto.fit(y, 100)
pareto_fitted = len(y)*pareto.pdf(range(artists.shape[0]), *param)
# param = gamma.fit(y) # gamma fits abysmally; see for yourself by uncomment
# gamma_fitted = len(y)*gamma.pdf(range(artists.shape[0]), *param)

pyplot.figure(figsize=(18, 6))
pyplot.bar(artists['Artist'], artists['Num Songs'])
pyplot.plot(pareto_fitted, color='r')
#pyplot.plot(gamma_fitted, color='g')
pyplot.xticks(visible=False)
pyplot.xlabel(artists.columns[0])
pyplot.ylabel(artists.columns[1])
pyplot.title('everybody');
```


```
/home/pavel/.local/lib/python3.6/site-packages/scipy/stats/_distn_infrastructure.py:2381: RuntimeWarning: invalid value encountered in double_scalars
    Lhat = muhat - Shat*mu
/home/pavel/.local/lib/python3.6/site-packages/scipy/stats/_distn_infrastructure.py:1682: RuntimeWarning: divide by zero encountered in log
    return log(self._pdf(x, *args))
```



Best fit is still too sharp for the data, and I tried for a good long while to get it to fit better, so I conclude this doesn't *quite* fit a power law.

Let's plot the top 50 artists so we can actually read who they are.

```
In [4]: pyplot.figure(figsize=(18, 10))
pyplot.bar(artists['Artist'][:50], artists['Num Songs'][:50])
pyplot.xticks(rotation=80)
pyplot.xlabel(artists.columns[0])
pyplot.ylabel(artists.columns[1])
pyplot.title('top 50');
```

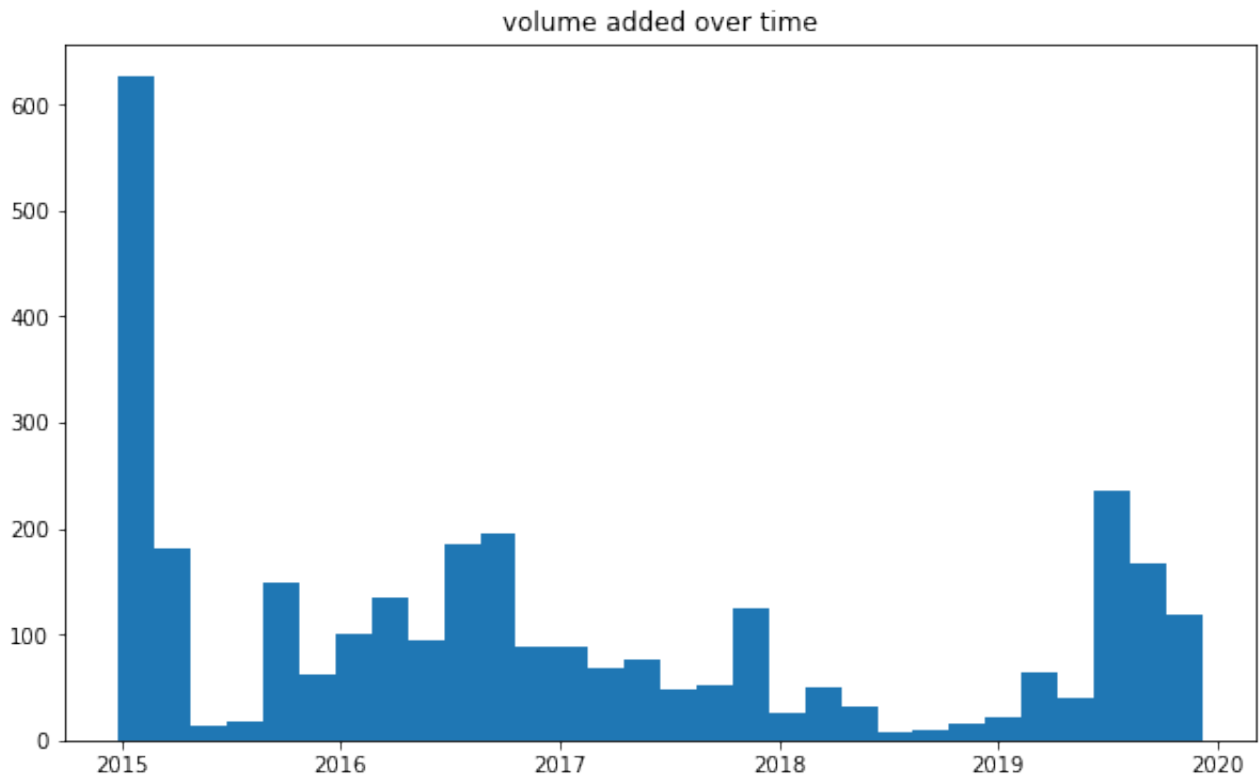


Volume Added Over Time

My proclivity to add songs to this playlist is a proxy for my interest in listening to music generally. How has it waxed and waned over time?

```
In [5]: from pandas.plotting import register_matplotlib_converters
register_matplotlib_converters() # to suppress warning

# Plot of added volume over time
parse_date = lambda d:(int(d[:4]), int(d[5:7]), int(d[8:10]))
pyplot.figure(figsize=(10, 6))
pyplot.hist([date(*parse_date(d)) for d in data['Added At']], bins=30)
pyplot.title('volume added over time');
```



The initial spike is from when I first started using Spotify as the home for this collection and manually added hundreds from my previous list.

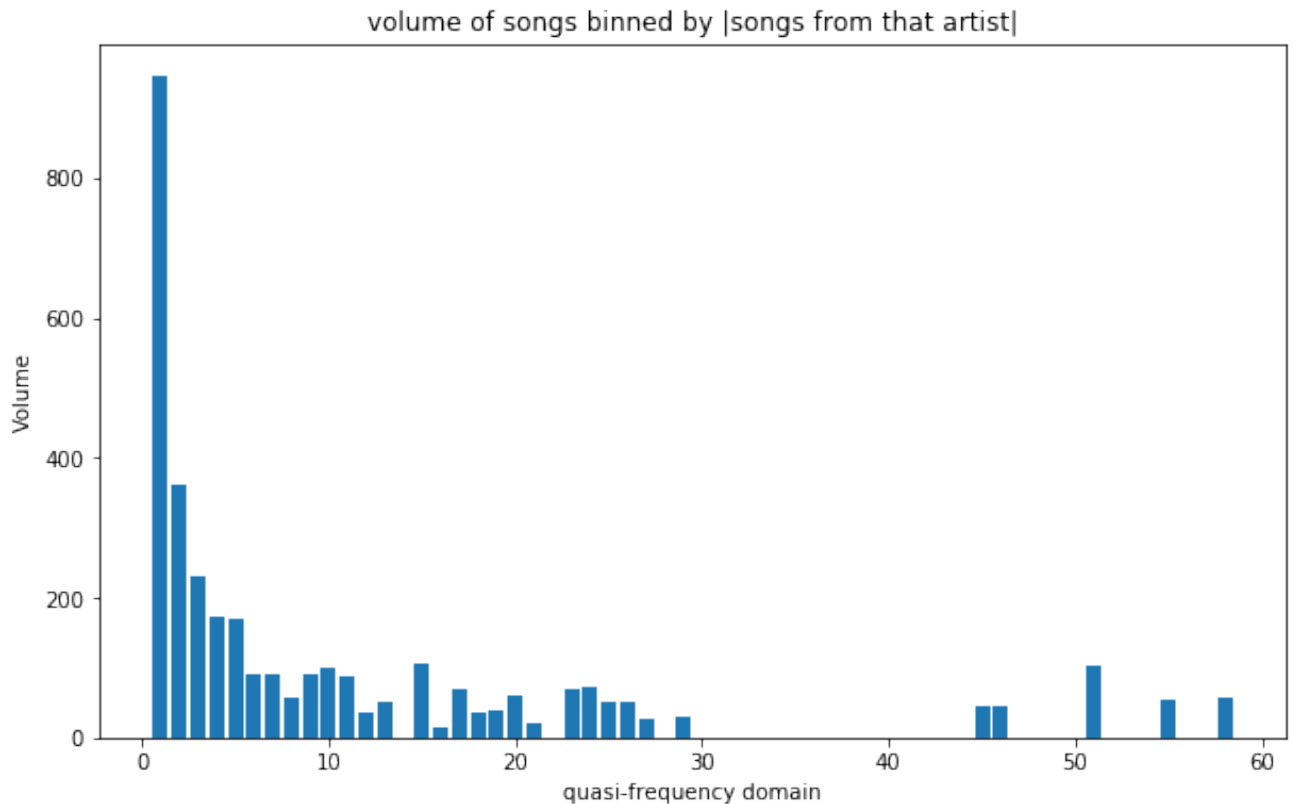
Eclecticism Measure (Frequency Transform)

This one is a personal favorite. I want to know how many of my songs are one-offs from that artist for me--just individual pieces I found fantastic and ended up adding after a few listens--, how many are two-offs, et cetera. I know it must be heavily skewed toward the low numbers.

```
In [6]: # bar chart of first bar chart == hipster diversity factor
frequency = defaultdict(int)
for n in artists['Num Songs']:
    frequency[n] += n
frequency = pandas.DataFrame(frequency.items(), columns=['Unique Count', 'Volume'])
frequency.sort_values('Volume', ascending=False)
print("number of song-artist pairs represented in the eclecticism chart:",
      sum(frequency['Volume']))

pyplot.figure(figsize=(10, 6))
pyplot.bar(frequency['Unique Count'].values, frequency['Volume'].values)
pyplot.title('volume of songs binned by |songs from that artist|')
pyplot.xlabel('quasi-frequency domain')
pyplot.ylabel(frequency.columns[1]);
```

number of song-artist pairs represented in the eclecticism chart: 3430



So, yes, it's much more common for an artist to make it in my list a few times than many times. In fact, the plurality of my top songs come from unique artists.

Conversely, this view also makes stark those few musicians from whom I've collected dozens.

Note that here, as in the artist bar charts, some songs are doubly-counted, because in cases artists collaborated I listed the song in both bins.

Genres Bar Chart

Alright, enough messing around. All the above were possible with the output from Watsonbox's Exportify. Let's get to the novel stuff you came *here* for.

People describe music by genre. As we'll see, genre names are flippin' hilarious and extremely varied, but in theory if I cluster around a few, that should give you a flavor of my tastes.

```
In [7]: # count songs per genre
genres = defaultdict(int)
for i,song in data.iterrows():
    if type(song['Genres']) is str: # some times there aren't any, and this
```

```


for genre in song['Genres'].split(','):
    if len(genre) > 0: # empty string seems to be a legit genre
        genres[genre] += 1

# sort for chart
genres = pandas.DataFrame(genres.items(), columns=['Genre', 'Num Songs']
                           ).sort_values('Num Songs', ascending=False).reset_
print("number of unique genres:", genres.shape[0])

pyplot.figure(figsize=(18, 6))
pyplot.bar(genres['Genre'], genres['Num Songs'])
pyplot.xticks(visible=False)
pyplot.xlabel(genres.columns[0])
pyplot.ylabel(genres.columns[1])
pyplot.title('All the genera');

```

number of unique genres: 755



So many! Let's do the same thing as with the artists and for giggles see if it fits a power law.

```

In [8]: y = []
for i in range(genres.shape[0]):
    for j in range(genres['Num Songs'][i]):
        y.append(i)

# sanity check
#pyplot.figure()
#pyplot.hist(y, bins=30)

param = pareto.fit(y, 100)
pareto_fitted = len(y)*pareto.pdf(range(genres.shape[0]), *param)


pyplot.figure(figsize=(18, 6))
pyplot.bar(genres['Genre'], genres['Num Songs'])
pyplot.plot(pareto_fitted, color='r')
pyplot.xticks(visible=False)

```



```
pyplot.xlabel(genres.columns[0])
pyplot.ylabel(genres.columns[1])
pyplot.title('All the genera');
```


```
/home/pavel/.local/lib/python3.6/site-packages/scipy/stats/_distn_infrastructure.py:2381: RuntimeWarning: invalid value encountered in double_scalars
  Lhat = muhat - Shat*mu
```



Still too sharp, but fits better than with the artists.

Let's look at the top 50 so we can read the names.

```
In [9]: pyplot.figure(figsize=(18, 10))
pyplot.bar(genres['Genre'][:50], genres['Num Songs'][:50])
pyplot.xticks(rotation=80)
pyplot.xlabel(genres.columns[0])
pyplot.ylabel(genres.columns[1])
pyplot.title('top 50');
```




"Indie popitism" lol. wtf? "Dreamo", "Vapor soul", "Freak folk", "Tropical house", "Post-grunge", "Hopebeat", "Noise pop", "Mellow gold"

These are too good. Next time someone asks me my music taste, I'm definitely using these.

If these are the *most* popular names, what are the really unique ones at the bottom of the chart?

```
In [10]: pyplot.figure(figsize=(18, 1))
pyplot.bar(genres['Genre'][-50:], genres['Num Songs'][-50:])
pyplot.xticks(rotation=80)
pyplot.xlabel(genres.columns[0])
pyplot.ylabel(genres.columns[1])
pyplot.title('bottom 50');
```



"hauntology", "psychedelic folk", "stomp and whittle", "dark trap", "filthstep",
 "shamanic", "deep underground hip hop", "future garage"

That was fun.

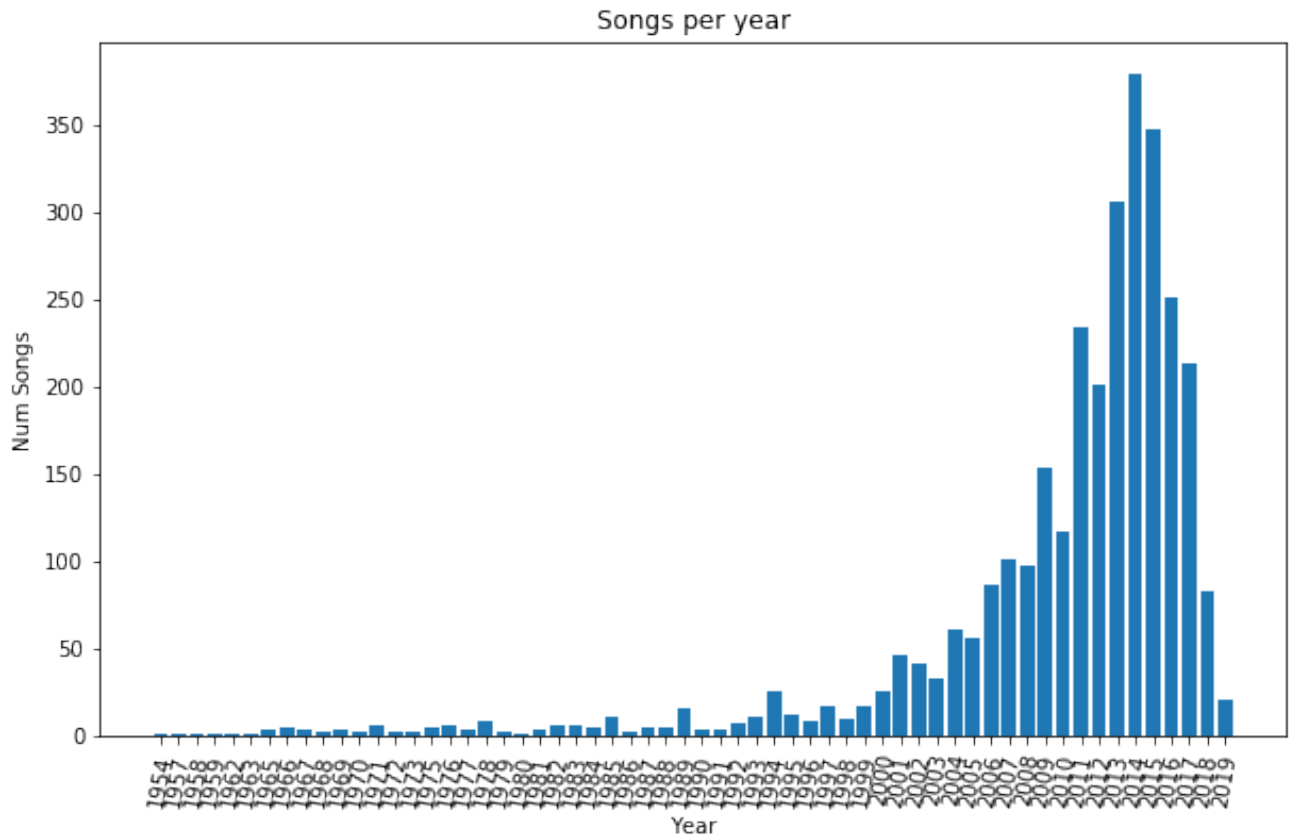
Release Dates

Which era of music do I prefer?

```
In [11]: years = defaultdict(int)
for i, song in data.iterrows():
    years[song['Release Date'][:4]] += 1

years = pandas.DataFrame(years.items(), columns=['Year', 'Num Songs'])
    .sort_values('Year')

pyplot.figure(figsize=(10, 6))
pyplot.bar(years['Year'], years['Num Songs'])
pyplot.xticks(rotation=80)
pyplot.xlabel(years.columns[0])
pyplot.ylabel(years.columns[1])
pyplot.title('Songs per year');
```



It seems to follow a Gamma distribution! This makes sense because I'm more likely to have heard things that are nearer me in time, and it takes a while for them to get through my process and become favorites.

Let's fit that gamma to the time-reversed data.

```
In [12]: # Some years are missing, so transform to a dataframe that covers full time
eldest = int(years['Year'].values[0])
youngest = int(years['Year'].values[-1])
missing_years = [str(x) for x in range(eldest+1, youngest) if
                  str(x) not in years['Year'].values]
ago = years.append(pandas.DataFrame.from_dict(
    {'Year': missing_years, 'Num Songs': [0 for x in range(len(missing_years)
    )].sort_values('Year', ascending=False).reset_index(drop=True)

y = []
for i in range(ago.shape[0]):
    for j in range(int(ago['Num Songs'][i])):
        y.append(i)

# sanity check histogram to make sure I'm constructing y properly
#pyplot.figure()
#pyplot.hist(y, bins=30)
```

```

param = gamma.fit(y, 10000)
gamma_fitted = len(y)*gamma.pdf(range(ago.shape[0]), *param)

pyplot.figure(figsize=(10, 6))
pyplot.bar(range(len(ago['Year'])), ago['Num Songs'])
pyplot.plot(gamma_fitted, color='g')
pyplot.xlabel('Years Ago')
pyplot.ylabel(ago.columns[1])
pyplot.title('Songs per year (in absolute time)');

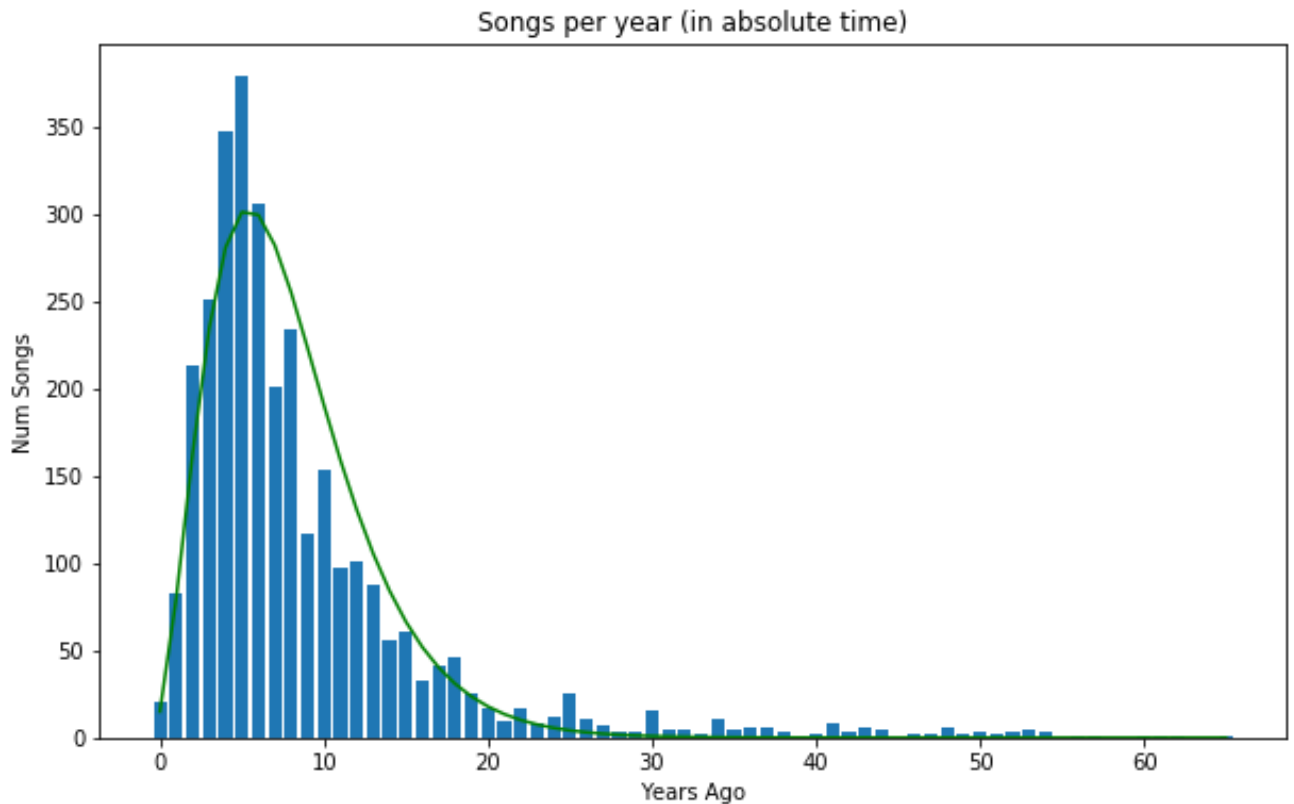
print('Oldest Hall of Fame')
print(data[['Track Name', 'Artist Name(s)', 'Release Date']].sort_values(
    'Release Date')[:10])

```

Oldest Hall of Fame

	Track Name \
3063	That's Amore
2509	(Where Do I Begin) Love Story
3021	Autumn Nocturne
2484	Take Five
2697	Stand by Me
0	Fanfare for the Common Man
554	Get Ready For This
1422	New Math
1905	Yesterday – Remastered 2009
2123	Il Buono, Il Brutto, Il Cattivo: Titoli Di Testa

	Artist Name(s)	Release Date
3063	Dean Martin,Dick Stabile And His Orchestra	1954
2509	Andy Williams	1957
3021	Lou Donaldson	1958
2484	The Dave Brubeck Quartet	1959-12-14
2697	Ben E. King	1962-08-20
0	Aaron Copland,London Symphony Orchestra	1963
554	2 Unlimited	1965
1422	Tom Lehrer	1965-01-01
1905	The Beatles	1965-08-06
2123	Ennio Morricone	1966



Pretty good fit! I seem to be extra partial to music from about 5 years ago. We'll see whether the present or maybe the further past catches up.

Popularity Contest

I was happy to find popularity listed as a field in Spotify's track JSON. It's a percentile between 0 and 100, rather than an absolute number of plays. Still, it can be used to give a notion of how hipster I am.

```
In [13]: popularity = defaultdict(int)
for i,song in data.iterrows():
    popularity[song['Popularity']] += 1

popularity = pandas.DataFrame(popularity.items(), columns=['Popularity', 'Num Songs'])
popularity.sort_values('Popularity')

pyplot.figure(figsize=(10, 6))
pyplot.bar(popularity['Popularity'].values, popularity['Num Songs'].values)
pyplot.xlabel(popularity.columns[0])
pyplot.ylabel(popularity.columns[1])
pyplot.title('popularity distribution');

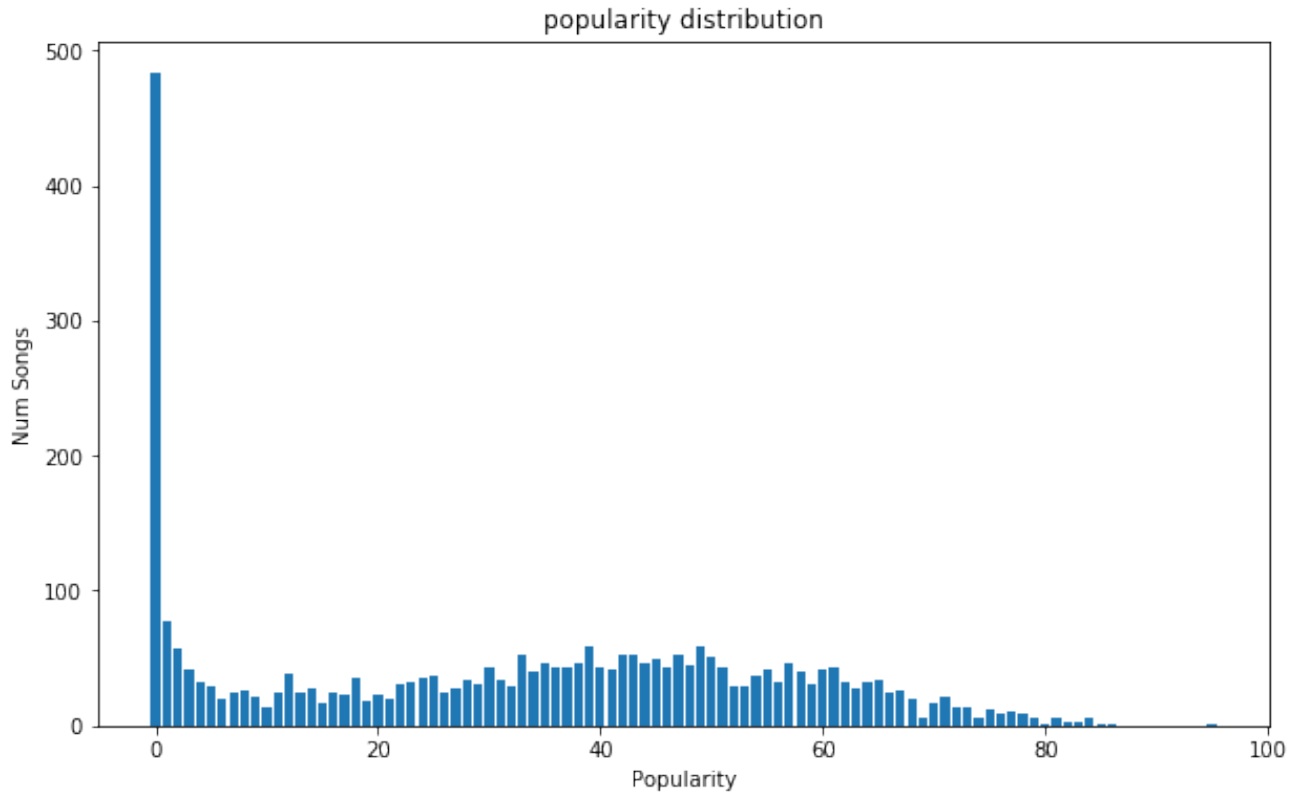
print("Average song popularity: ", popularity['Popularity'].mean())
print("Median song popularity: ", popularity['Popularity'].median())
```

```
print("Max song popularity: ", popularity['Popularity'].max())
```

Average song popularity: 43.59090909090909

Median song popularity: 43.5

Max song popularity: 95



Damn, I'm a hipster.

Track Duration

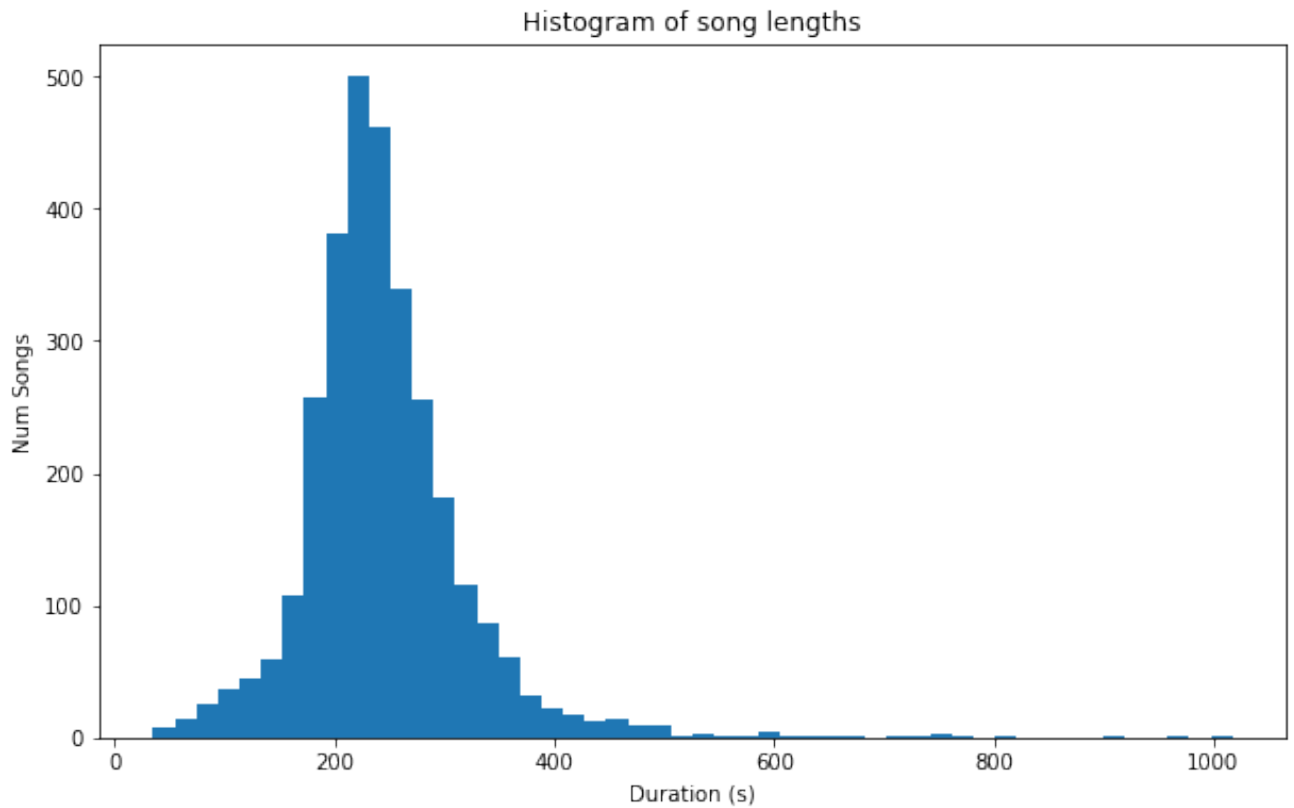
Do I prefer long songs or short ones?

```
In [14]: pyplot.figure(figsize=(10,6))
pyplot.hist(data['Duration (ms)']/1000, bins=50);
pyplot.xlabel('Duration (s)')
pyplot.ylabel('Num Songs')
pyplot.title('Histogram of song lengths')

mean = data['Duration (ms)'].mean()/1000
median = data['Duration (ms)'].median()/1000
print("Average song length: " + str(int(mean//60)) + (":" if mean%60 >=10 else
+ str(mean%60))
print("Median song length: " + str(int(median//60)) + (":" if median%60 >=10
+ str(median%60))
```

Average song length: 4:06.183588730569966

Median song length: 3:56.16



Median is lower than the mean, so I'm skewed right. That is, I like a few really long songs. What are they?

```
In [15]: print("Longest Hall of Fame:")
print(data[['Track Name', 'Artist Name(s)', 'Release Date', 'Duration (ms)']
          'Duration (ms)', ascending=False)[:10])
```


Longest Hall of Fame:

	Track Name \
705	Irene
1954	The Return of the King (From The Lord of the R...
464	The Cure For Pain
2406	Shine On You Crazy Diamond (Pts. 1-5)
142	Two Step - Live At Piedmont Park, Atlanta, GA ...
1474	Cage-Nerd
2407	Shine On You Crazy Diamond (Pts. 6-9)
144	Warehouse - Live At Piedmont Park, Atlanta, GA...
143	Don't Drink the Water - Live At Piedmont Park,...
2717	The Alien

	Artist Name(s)	Release Date	Duration (ms)
705	Beach House	2012-05-15	1017013
1954	The City of Prague Philharmonic Orchestra	2004-01-01	976893
464	mewithoutYou	2002-01-01	908840
2406	Pink Floyd	1975-09-12	811077
142	Dave Matthews Band	2007-12-11	808226
1474	Tim Minchin	2011-04-04	778250
2407	Pink Floyd	1975-09-12	747325
144	Dave Matthews Band	2007-12-11	743906
143	Dave Matthews Band	2007-12-11	743493
2717	Ben Salisbury, Geoff Barrow	2018-02-23	723579

Musical Features

In the interest of understanding user tastes and providing the best possible music recommendations, Spotify has done [some really sophisticated analysis](#) of actual track content. Music is a time series, but most similarity metrics (and most ML methods generally) require inputs to be vectors, that is: points in some feature-space. So they've transformed the tracks to numerical metrics like Energy and Valence (continuous) and Key (discrete).

For the continuous metrics, they [provide distributions across all music](#). Here they are next to similar plots of my own songs.

```
In [16]: pyplot.figure(figsize=(20,40))

for i,category in enumerate(['Tempo', 'Acousticness', 'Instrumentalness', 'L
                             'Valence', 'Speechiness', 'Loudness', 'Energy',
pyplot.subplot(9, 2, i*2+1)
# It would be nice to show the KDE on these plots, but there isn't a way
# to show it on unnormalized https://github.com/mwaskom/seaborn/issues/4
pyplot.hist(data[category], bins=30)
pyplot.text(min(data[category]), 0, r'$\mu$'+str(data[category].mean()))
pyplot.xlabel('Value')
```

```











pyplot.ylabel('Num Songs')
pyplot.title(category)

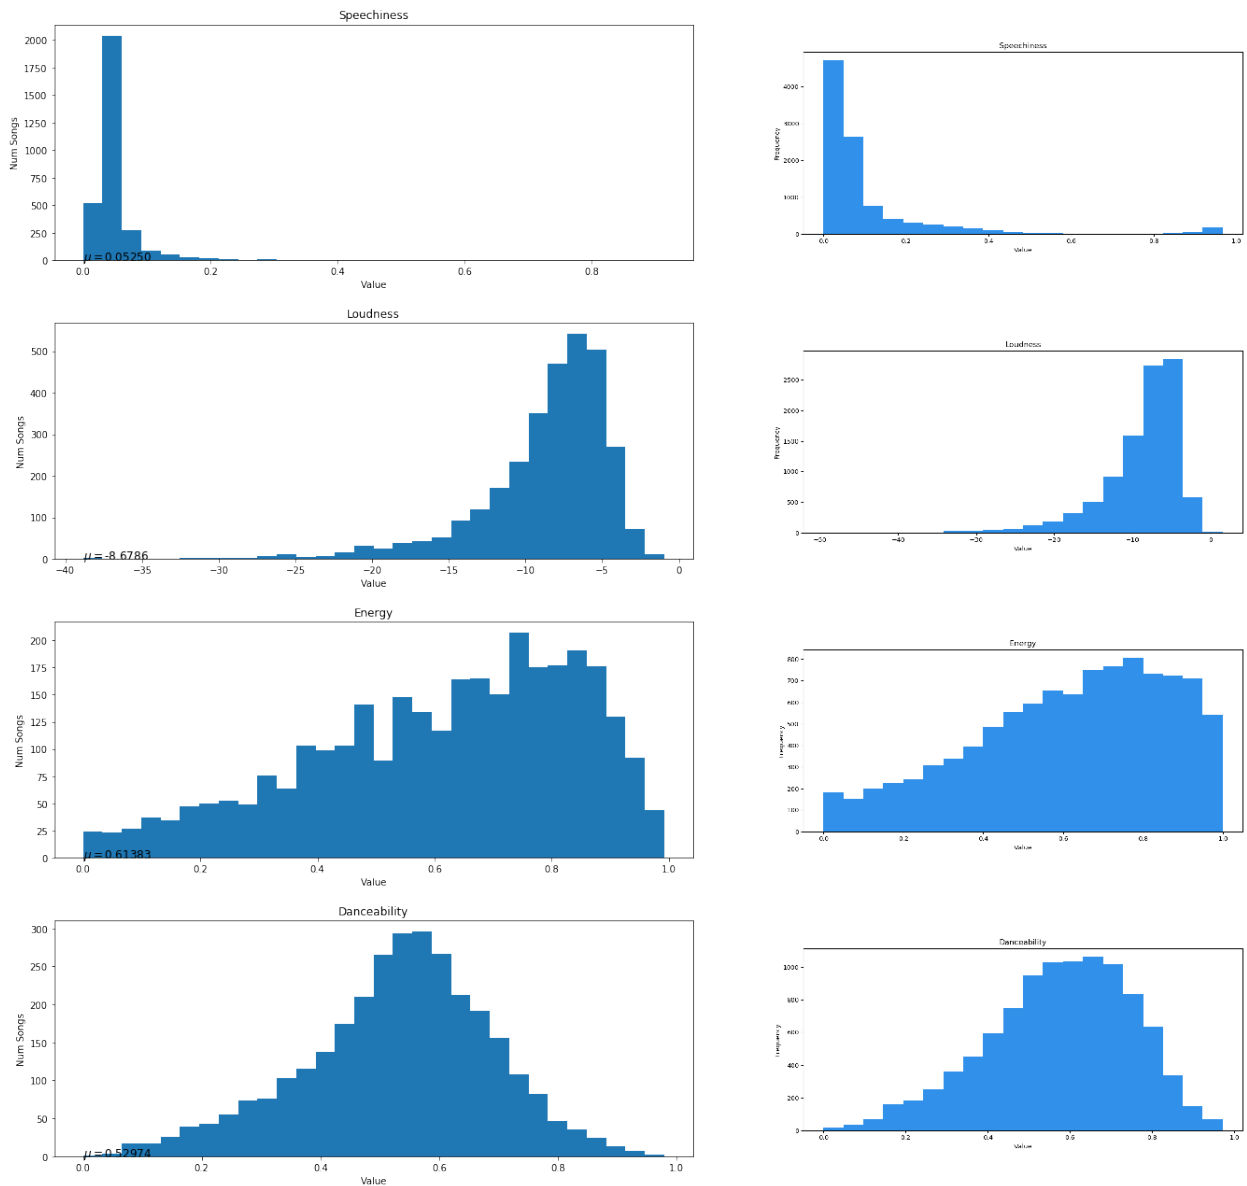
pyplot.subplot(9, 2, i*2+2)
pyplot.imshow(pyplot.imread('https://developer.spotify.com/assets/audio/
category.lower()+'.png'))

pyplot.axis('off')

pyplot.tight_layout(h_pad=2)

```





Looks like my preferred Tempo, Acousticness, Instrumentalness, Liveness, Speechiness, and Loudness are not much different from average. Energy is pretty similar, but I have perhaps slightly lower affinity for the super-energetic stuff. My Valence is somewhat negatively skewed, meaning I like sadder songs than average. And my Danceability peaks lower than average.

Now let's look at the discrete music features.

```
In [17]: pyplot.figure(figsize=(15,4))

pyplot.subplot(1, 3, 1)
seaborn.countplot(data['Time Signature'])
pyplot.xlabel('Beats per bar')
pyplot.ylabel('Num Songs')
```

```


pyplot.title('Time Signature')

pyplot.subplot(1, 3, 2)
axes = seaborn.countplot(data['Key'])
axes.set(xticklabels=['C', 'C#', 'D', 'D#', 'E', 'F', 'F#', 'G', 'G#', 'A',
pyplot.ylabel('Num Songs')
pyplot.title('Key')

pyplot.subplot(1, 3, 3)
axes = seaborn.countplot(data['Mode'])
axes.set(xticklabels=['minor', 'major'])
pyplot.ylabel('Num Songs')
pyplot.title('Major vs Minor Key');

pyplot.tight_layout(w_pad=2)

```



Musicians seem to favor C major and eschew D#. More than a third of my songs are in a minor key. I don't have a baseline to compare against here, but this might contribute to my lower Valence.

Looks like the vast majority of my music is 4/4 time with a good few in 3/4. I wasn't even aware there were any with 5 beats. What are those?

```

In [18]: print('5:\n', data.loc[data['Time Signature']==5][
          ['Track Name', 'Artist Name(s)', 'Release Date'][:20])

```

5:

	Track Name \
76	Yachts – A Man Called Adam mix
121	Good Morning Fire Eater
227	Carry On
248	Elysium
277	Lately
390	Evenstar
451	Make A Fist
463	(B)
573	Animals
741	All That Remains
749	Crush The Camera
827	0
1081	Cold Sparks
1182	You Are Gonna Die
1209	Everything In Its Right Place
1214	The Tourist
1215	I Am Citizen Insane
1875	Have I Always Loved You?
2015	Resonance
2180	Pray

	Artist Name(s)	Release Date
76	Coco Steel Lovebomb	2000-10-31
121	Copeland	2008-01-01
227	fun.	2012-02-21
248	Klaus Badelt,Lisa Gerrard,Gavin Greenaway,The ...	2000-01-01
277	Memoryhouse	2011-09-13
390	Howard Shore,Isabel Bayrakdarian	2002-12-02
451	Phantogram	2011
463	mewithoutYou	2002-01-01
573	Muse	2012-09-24
741	Rogue Wave	2010
749	Rogue Wave	2005-08-23
827	Coldplay	2014-05-19
1081	Mutemath	2011-09-30
1182	Marc Streitenfeld	2015-03-24
1209	Radiohead	2000-10-02
1214	Radiohead	1997-06-17
1215	Radiohead	2003-06-09
1875	Copeland	2014-11-17
2015	Home	2014-07-01
2180	Sam Smith	2017-10-06

Make A Fist is totally 5/4, and so is Animals. Funny how I didn't notice the strange energetic time signature until now. But Carry On is definitely 4/4, as is Yachts, and Pray is 6/8. So Spotify's algorithm isn't perfect at this, which is expected.

What are 0 and 1?

```
In [19]: print('0:\n', data.loc[data['Time Signature']==0][
        ['Track Name', 'Artist Name(s)', 'Release Date'][:10])
        print('\n1:\n', data.loc[data['Time Signature']==1][
        ['Track Name', 'Artist Name(s)', 'Release Date'][:20])
```

0:

	Track Name	Artist Name(s)	Release Date
1393	Small Memory	Jon Hopkins	2009-05-05

1:

	Track Name \
71	Clair De Lune
120	Top Of The Hill
231	I Am the Very Model of a Modern
243	The Last of Us (You and Me)
366	Bowery
507	The Eternal City
570	Prelude
608	Pú ert jörðin
611	Raein
1302	Campfire Song Song
1356	Mylo Xyloto
1399	Anagram
1957	The Fellowship (From The Lord of the Rings: Th...
1997	Monsoon
2041	Meet Me in the Woods
2080	Only Songs
2228	Old Casino
2241	Work This Time
2658	I Don't Think So
2738	Other Worlds

	Artist Name(s)	Release Date
71	Claude Debussy	2014-10-13
120	Conduits	2013-04-16
231	The Pirates Of Penzance	1981
243	Gustavo Santaolalla, Alan Umstead	2013-06-07
366	Local Natives	2013-01-29
507	Michele McLaughlin	2007-12-04
570	Muse	2012-09-24
608	Ólafur Arnalds	2010-05-07
611	Ólafur Arnalds	2009-08-28
1302	Spongebob Squarepants	2001
1356	Coldplay	2011-10-24
1399	Young the Giant	2014-01-17
1957	The City of Prague Philharmonic Orchestra	2004-01-01
1997	Hippo Campus	2017-02-24
2041	Lord Huron	2015-04-07
2080	The Wild Reeds	2017-04-07
2228	Coastgaard	2016-02-26
2241	King Gizzard & The Lizard Wizard	2014-03-07
2658	Ben Phipps	2016-09-30
2738	Bassnectar, Dorfex Bos	2017-12-01

Looks like there is only one song with 0 time signature. It's a piano piece with a tempo

that rises and falls. This category might be for variable tempo.

Claire De Lune is 9/8 time, so sort of waltzish but not really.

The Major General's Song is 4/4, but there are some stops in there and a lot of speaking, so I understand how that might be difficult to pick out.

Top of the Hill really sounds like 7/4 to me (1-2-123 sort of beat).

Þú ert jörðin is actually properly 1/4 time according to the internet, and relistening I understand how that could be the case. It's like there are little riffs each bar following a quadruplet pattern, but the major beats really only come every bar.

The Last of Us (You and Me) seems similar. It might be properly 1/4 time.

So it looks like this category is for actual single beats and unusual time signatures that Spotify isn't sure what to do with.

Joint Analysis


I mostly just want to showcase what's possible. Let's plot Energy and Popularity together to see whether there is a relationship.

```
In [20]: x = 'Energy'
y = 'Popularity'

axes = seaborn.jointplot(x=data[x], y=data[y], kind='hex', color='r', size=1
axes.set_axis_labels(x, y, fontsize=20);
```

```
/home/pavel/.local/lib/python3.6/site-packages/seaborn/axisgrid.py:2262: Use
rWarning: The `size` paramter has been renamed to `height`; please update yo
ur code.
```

```
warnings.warn(msg, UserWarning)
```

The data is pretty scattered around the whole plot, meaning the relationship here is actually pretty weak. Surprising.

The Final Frontier

Finally, I'm going to follow [this guy's example](#) and feed the dimension-reduced data to a one-class SVM to get a sense of what the frontier of my normal taste looks like in that space, [heat-map-of-the-universe-style](#).

[t-SNE](#) is a method for visualizing high-dimensional data in low-dimension. Songs which are more alike will be nearer each other in the feature space, but we can't visualize a

space with that many dimensions. What we can do is reconstitute the points in 2D, attempting to preserve the pairwise distances, the notions of similarity, between songs.

```
In [21]: show_percent = 2

from sklearn.manifold import TSNE
from random import random
from sklearn.svm import OneClassSVM
import numpy

# Create a dataframe of only the numerical features, all normalized so embed
# doesn't get confused by scale differences
numerical_data = data.drop(['Spotify ID', 'Artist IDs', 'Track Name',
                            'Album Name', 'Artist Name(s)', 'Added By', 'Added At',
                            'Genres'], axis=1)
numerical_data['Release Date'] = pandas.to_numeric(
    numerical_data['Release Date'].str.slice(0,4))
numerical_data = (numerical_data - numerical_data.mean())/numerical_data.sta
print('using:', list(numerical_data.columns))

# If you like, only include a subset of these, because the results with all
# is really hard to interpret
#tsne_data = numerical_data[['Popularity', 'Energy', 'Acousticness',
#                             'Valence', 'Loudness']]
#
#print("\nConsidering similarity with respect to the following features:")
#print(tsne_data.dtypes)

# Takes a 2D data embedding and an svm trained on it and plots the decision
def plotFrontier(embedded, svm, technique_name, scale):
    # get all the points in the space, and query the svm on them
    xx, yy = numpy.meshgrid(numpy.linspace(min(embedded[:,0])*scale,
                                              max(embedded[:,0])*scale, 500),
                            numpy.linspace(min(embedded[:,1])*scale,
                                              max(embedded[:,1])*scale, 500))

    Z = svm.decision_function(numpy.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape) # positive Z means yes. negative means outliers.

    pyplot.figure(figsize=(20,20))
    pyplot.title('Decision boundary of One-class SVM in '+technique_name+' s
    pyplot.contourf(xx, yy, Z, levels=numpy.linspace(Z.min(), 0, 7), cmap=py
    pyplot.contour(xx, yy, Z, levels=[0], linewidths=2, colors='green') # th
    pyplot.contourf(xx, yy, Z, levels=[0, Z.max()], colors='lightgreen')

    pyplot.scatter(embedded[:, 0], embedded[:, 1], s=10, c='grey')
    for i,song in data.iterrows():
        if random() < show_percent*0.01: # randomly label % of points
            #if song['Artist Name(s)'] in ['Coldplay']:
                x, y = embedded[i]
                pyplot.annotate(song['Track Name'], (x,y), size=10,
```

```

xytext=(-30,30), textcoords='offset points',
ha='center',va='bottom',
arrowprops={'arrowstyle':'->', 'color':'red'})

tsne_embedded = TSNE(n_components=2).fit_transform(numerical_data)

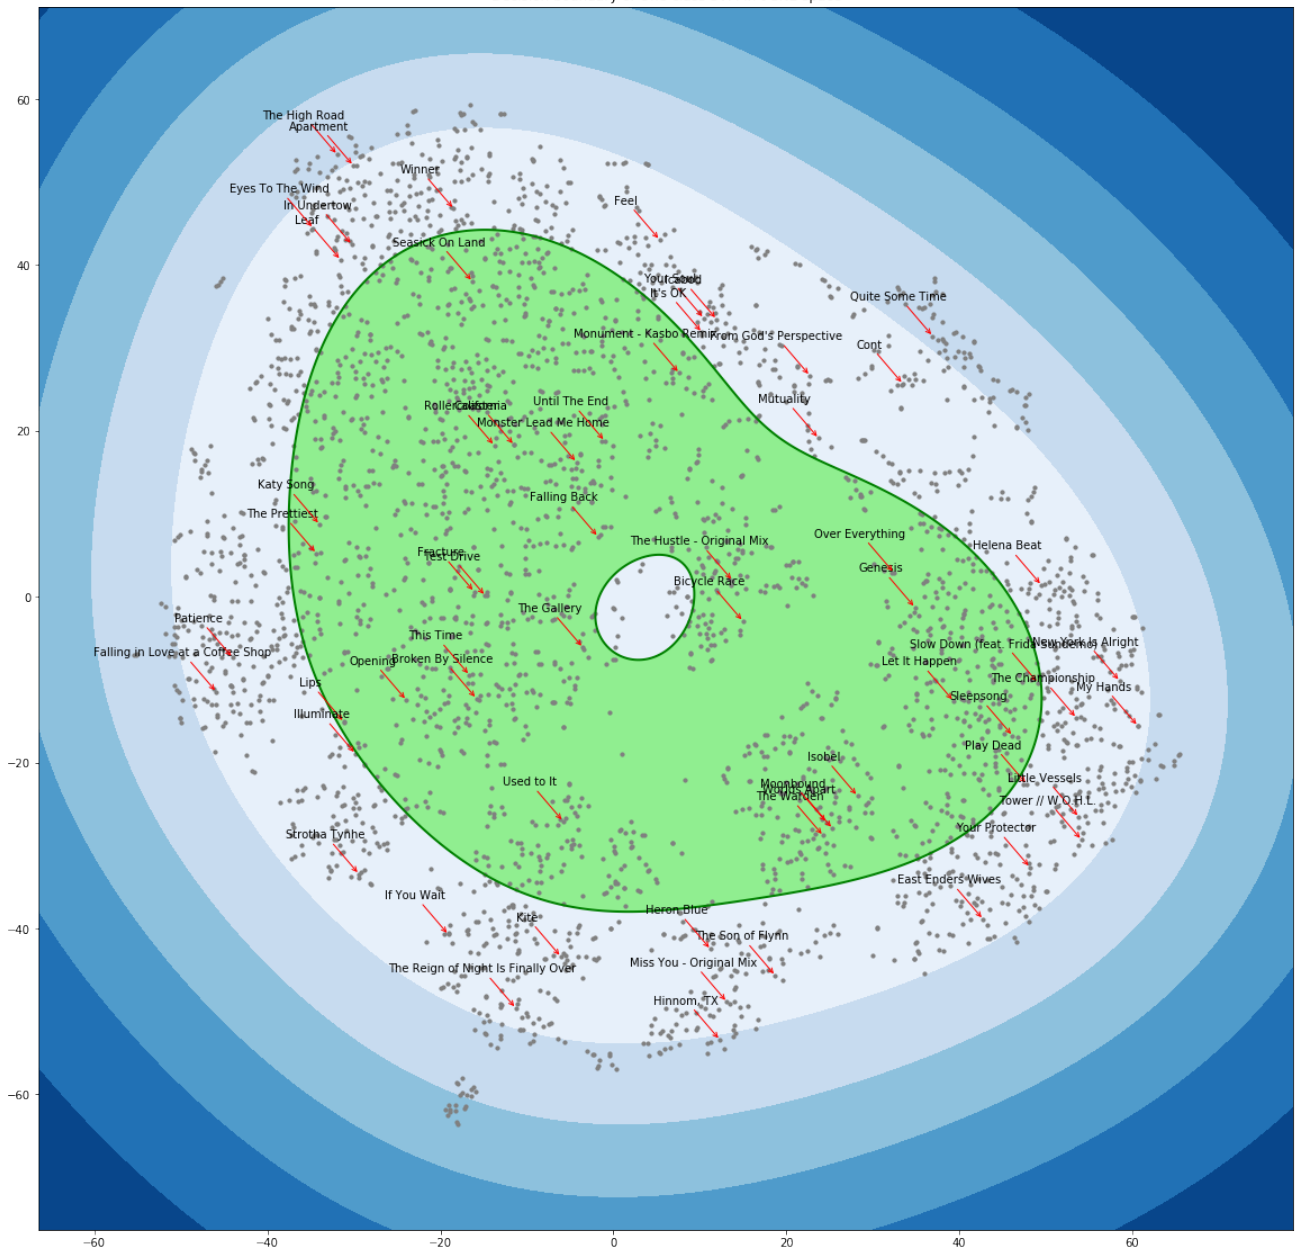
svm_tsne = OneClassSVM(gamma='scale')
svm_tsne.fit(tsne_embedded)

plotFrontier(tsne_embedded, svm_tsne, 't-SNE', 1.2)

```

using: ['Release Date', 'Duration (ms)', 'Popularity', 'Danceability', 'Energy', 'Key', 'Loudness', 'Mode', 'Speechiness', 'Acousticness', 'Instrumentalness', 'Liveness', 'Valence', 'Tempo', 'Time Signature']

Decision boundary of One-class SVM in t-SNE space



The point scatter looks really different every time this runs, because it's stochastic. The clusters don't necessarily have sensible interpretations, though you might be able to label a few of them. It's good to see some notionally similar pieces ending up near each other. You can try this with a subset of these dimensions to try to make the result more interpretable.

Modifying the parameters of the SVM changes its fit significantly, so I'm not sure this is the best model. Gamma too large just clearly overfits the data. Gamma too small just makes the decision boundary a boring ellipse. Using `gamma='scale'` as [the docs](#) recommend is a more interesting middle ground, but still the SVM seems to believe that a great many of the songs I love fall outside the boundary.

I'll try a different dimensionality reduction technique. [The original author](#) uses Principle Component Analysis to feed the SVM.

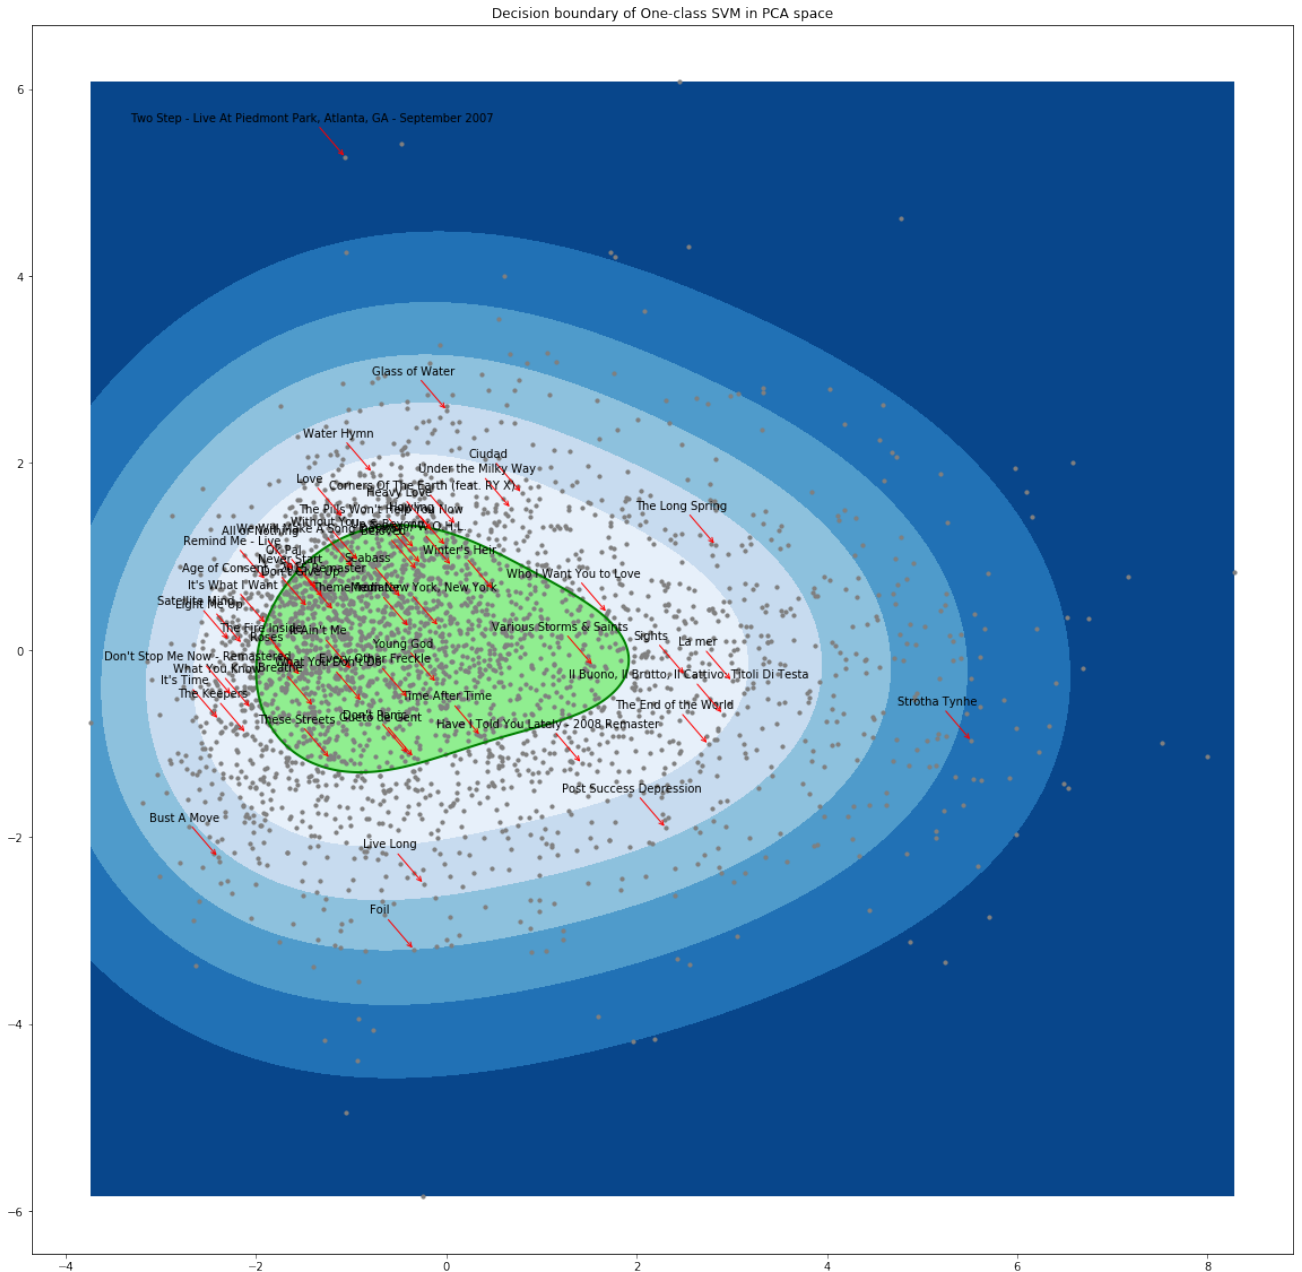
```
In [22]: from sklearn.decomposition import PCA
```

```
pca = PCA(n_components=2)
pca_embedded = pca.fit_transform(numerical_data)
print("% variance explained by successive PCA dimensions:",
      pca.explained_variance_ratio_)

svm_pca = OneClassSVM(gamma='scale')
svm_pca.fit(pca_embedded)

plotFrontier(pca_embedded, svm_pca, 'PCA', 1)
```

```
% variance explained by successive PCA dimensions: [0.21986282 0.09180652]
```



Ideally, songs falling nearer the center here, like Cheeseburger in Paradise and RAC's We Belong, are those that most characterize my taste numerically, and the odd ones, like Pink Floyd's Comfortably Numb and The Fellowship of the Ring orchestral suite, fall on the outside.

So in the end my music taste is a blob that doesn't even fit the data very well. And that's the point: Like many things, it's too complicated to boil down. You can't answer the question fully. But understanding elements of the answer can aid the process of discovery, and that's valuable. It's why Spotify is such a force at music recommendation. It's why Data Science.